



Deep Learning with Shallow Pixels: Classifying Images

A Transfer Learning Approach


Project for Masterschool
by Dido De Boedt



About the project

Goal

Build an image classifier
for the 10 classes of
CIFAR-10



Approach

Transfer learning using
ResNet50

pre-trained on ImageNet

Limited to 10,000 training
samples for faster
experimentation

Tools

TensorFlow (Keras API in
Google Colab)

About the data




- **CIFAR-10** is a computer vision dataset
- Contains 60,000 32x32 color images
- Images are RGB with rich variation in style and context
- 10 mutually exclusive Classes

ResNet50 Model

- 50-layer CNN with residual connections
- Solves vanishing gradient problem by skipping layers
- Pretrained on ImageNet (1M+ images, 1000 classes)
- Excellent for transfer learning to new tasks like CIFAR-10

Residual connections act like Lego tunnels that skip layers, helping information flow through deep networks.



Model Architecture


Base model:

- ResNet50 pre-trained on ImageNet
- Input resized to 32x32

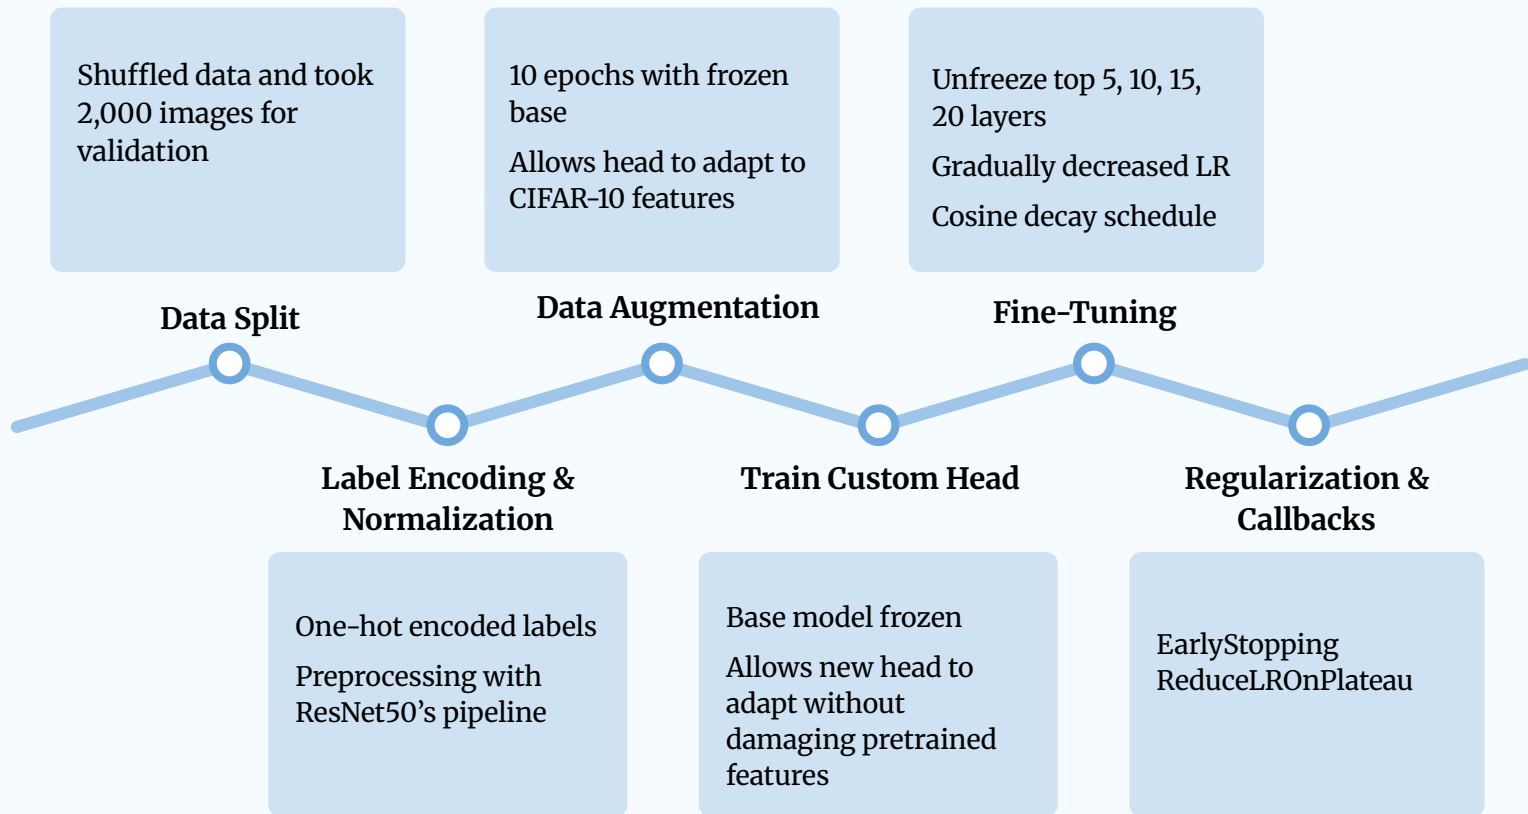
Custom classification head:

- Dense layers with ReLU
- BatchNormalization
- Dropout for regularization
- Softmax for 10 classes

Label smoothing: Encourages generalization



Preprocessing & Training



Model Evaluation & Deployment

About This Project

Navigation

- Go to:
- Upload Your Image
 - CIFAR-10 Samples
 - Analytics Dashboard
 - Model Insights

Session Stats

Total Predictions: 54
Session ID: 07e82f46

Clear Feedback Logs

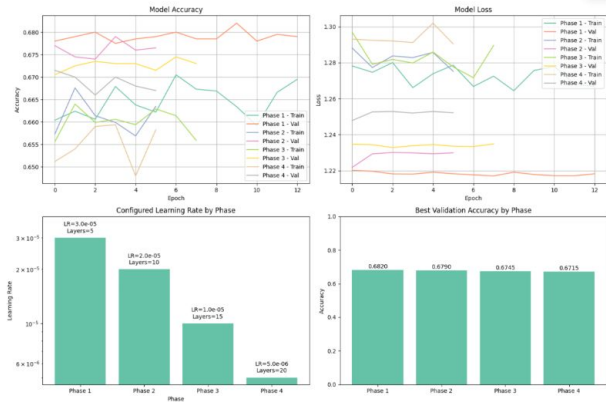
Advanced Settings

CIFAR-10 AI Image Classifier


Powered by ResNet50 Deep Learning | Computer Vision Project by Dido De Boott

Model Training Visualizations

Training Metrics



Confusion Matrix (Test Set 10,000 images)



Challenges

Small Training Set (10K)



Low Resolution (32x32 - ResNet50 expects 224x224)

Confusing Classes: Cats, dogs, birds look similar

Longer Training Times: Hyperparameter tuning needed patience

Learning Rate Scheduling:

- Needed careful adjustment to avoid divergence
- Smaller rates required for deeper unfreezing



Conclusion

- ✓ Transfer learning works even with limited data
- ✓ Layer-wise unfreezing + tuned learning rates
- ✓ Visualization crucial for diagnosing training
- ✓ Data augmentation essential for small datasets

Analytics Dashboard

Overview Metrics

Total Feedback

53

Accuracy

75.0%

Avg Confidence


62.1%

Unique Sessions

1

Accuracy by Class

Prediction Accuracy by Class (CIFAR Samples)





Thank You for your attention!

Project for Masterschool
by Dido De Boodt

